

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv(r"C:\Users\rubin\Downloads\bottle.csv")
df
```

C:\Users\rubin\AppData\Local\Temp\ipykernel_4564\2619180729.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.

```
df=pd.read_csv(r"C:\Users\rubin\Downloads\bottle.csv")
```

Out[2]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	I
				19-4903CR-HY-060-0930-05400560-0000A-3								
0	1	1	054.0 056.0		0	10.500	33.4400	NaN	25.64900	NaN	...	
				19-4903CR-HY-060-0930-05400560-0008A-3								
1	1	2	054.0 056.0		8	10.460	33.4400	NaN	25.65600	NaN	...	
				19-4903CR-HY-060-0930-05400560-0010A-7								
2	1	3	054.0 056.0		10	10.460	33.4370	NaN	25.65400	NaN	...	
				19-4903CR-HY-060-0930-05400560-0019A-3								
3	1	4	054.0 056.0		19	10.450	33.4200	NaN	25.64300	NaN	...	
				19-4903CR-HY-060-0930-05400560-0020A-7								
4	1	5	054.0 056.0		20	10.450	33.4210	NaN	25.64300	NaN	...	
...	
				20-1611SR-MX-310-2239-09340264-0000A-7								
864858	34404	864859	093.4 026.4		0	18.744	33.4083	5.805	23.87055	108.74	...	
				20-1611SR-MX-310-2239-09340264-0002A-3								
864859	34404	864860	093.4 026.4		2	18.744	33.4083	5.805	23.87072	108.74	...	
				20-1611SR-MX-310-2239-09340264-0005A-3								
864860	34404	864861	093.4 026.4		5	18.692	33.4150	5.796	23.88911	108.46	...	
				20-1611SR-MX-310-2239-09340264-0010A-3								
864861	34404	864862	093.4 026.4		10	18.161	33.4062	5.816	24.01426	107.74	...	

Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	I
864862	34404	864863	093.4	20-1611SR-	15	17.533	33.3880	5.774	24.15297	105.66	...
			026.4	MX-310-							
				2239-							
				09340264-0015A-3							

864863 rows × 74 columns

```
In [3]: df=df[['Salnty','T_degC']]
df.columns=['Sal','Temp']
```


```
In [4]: df.head()
```

Out[4]:

	Sal	Temp
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45

```
In [5]: sns.lmplot(x='Sal',y='Temp',data=df,order=2,ci=None)
```

```
Out[5]: <seaborn.axisgrid.FacetGrid at 0x11e503f6170>
```



```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 864863 entries, 0 to 864862  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0    Sal      817509 non-null  float64  
1    Temp     853900 non-null  float64  
dtypes: float64(2)  
memory usage: 13.2 MB
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	Sal	Temp
count	817509.000000	853900.000000
mean	33.840350	10.799677
std	0.461843	4.243825
min	28.431000	1.440000
25%	33.488000	7.680000
50%	33.863000	10.060000
75%	34.196900	13.880000
max	37.034000	31.140000

```
In [8]: df.fillna(method='ffill')
```

```
Out[8]:
```

	Sal	Temp
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

```
In [9]: df.fillna(value=0,inplace=True)
```

C:\Users\rubin\AppData\Local\Temp\ipykernel_4564\1434098079.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(value=0,inplace=True)
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: Sal      0  
Temp      0  
dtype: int64
```

```
In [11]: x=np.array(df['Sal']).reshape(-1,1)  
y=np.array(df['Temp']).reshape(-1,1)
```

```
In [12]: df.isna().any()
```

```
Out[12]: Sal      False  
Temp      False  
dtype: bool
```

```
In [13]: df.dropna(inplace=True)
```

C:\Users\rubin\AppData\Local\Temp\ipykernel_4564\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

```
In [14]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
reg=LinearRegression()  
reg.fit(x_train,y_train)  
print(reg.score(x_test,y_test))
```


```
0.014074449093311903
```

```
In [16]: y_pred=reg.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```





```
In [17]: df500=df[:][:500]  
sns.lmplot(x='Sal',y='Temp',data=df500,order=1,ci=None)
```

```
Out[17]: <seaborn.axisgrid.FacetGrid at 0x11e51498eb0>
```



```
In [20]: df500.fillna(method='ffill',inplace=True)
x=np.array(df500['Sal']).reshape(-1,1)
y=np.array(df500['Temp']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
reg=LinearRegression()
reg.fit(x_train,y_train)
print("Regression:",reg.score(x_test,y_test))
y_pred=reg.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```

Regression: 0.061335622373488774



```
In [21]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [22]: model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```


R2 score: 0.061335622373488774

IMPLEMENTING RIDGE AND LASSO REGRESSION

```
In [23]: from sklearn.linear_model import Ridge  
from sklearn.linear_model import RidgeCV  
from sklearn.linear_model import Lasso
```

```
In [24]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[24]: <Axes: >



```
In [25]: features = df.columns[0:2]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X_train is (605404, 2)
The dimension of X_test is (259459, 2)

```
In [26]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

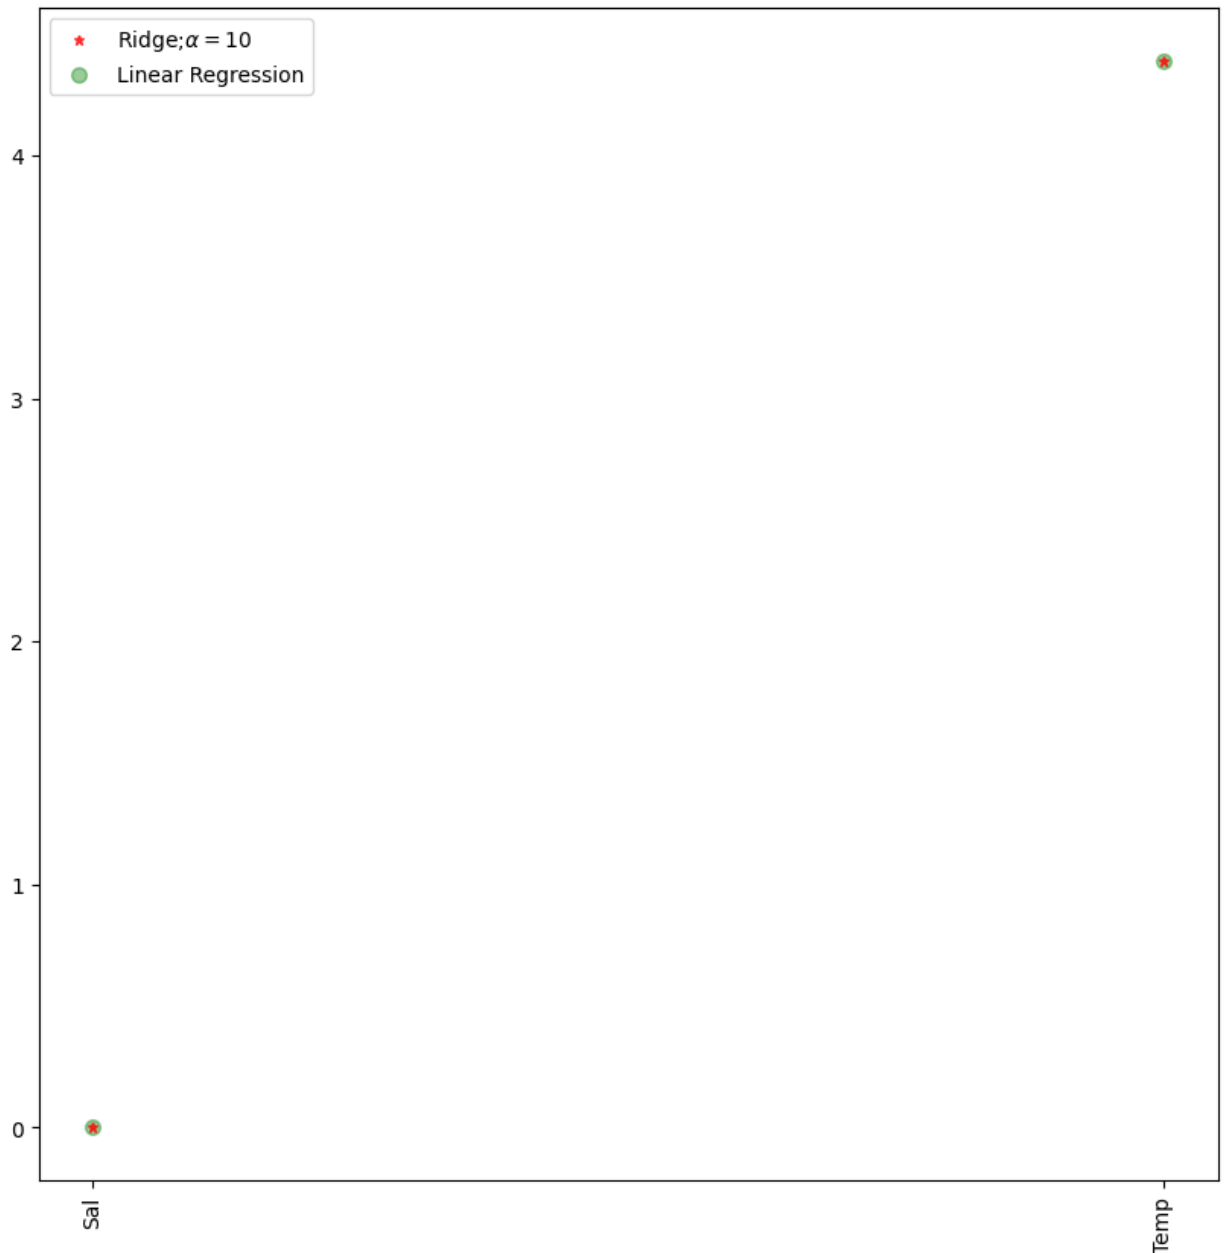
The train score for lr model is 1.0
The test score for lr model is 1.0

```
In [27]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.999999999723243
The test score for ridge model is 0.9999999997231402

```
In [28]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
#plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color=
plt.xticks(rotation=90)
plt.legend()
#plt.title("comparison plot of Ridge,Lasso and Linear regression model")
plt.show()
```



```
In [29]: print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -1.9031696447013857e-05

```
In [30]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[30]: <Axes: >




```
In [35]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train, y_train)
print(lasso_cv.score(X_train,y_train))
print(lasso_cv.score(X_test,y_test))
```

0.9999999994806811

0.9999999994806712

In [36]:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,
plt.plot(lasso_cv.coef_,alpha=0.6,linestyle='none',marker='d',markersize=6,color='bl
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color=
plt.xticks(rotation=90)
plt.legend()
plt.title("comparison plot of Ridge,Lasso and Linear regression model")
plt.show()
```



```
In [37]: from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test))
```

The train score for ridge model is 0.9999999981135502

The train score for ridge model is 0.99999999811206

ELASTIC NET REGRESSION

```
In [38]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(X,y)
print(regr.coef_)
print(regr.intercept_)
```

[0. 0.94934511]
0.5401219631063316

```
In [39]: y_pred_elastic=regr.predict(X_train)
```

```
In [40]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 114.40984808660129

```
In [ ]:
```